

# CS6650 Assignment2 Report

Xianhe Zhang

ZHANG.XIANH@

## Grading Note

I have talked about late submission with Ian. The below is communication email.

**Re: Request about CS6650 HW2 DDL extension.**



 **Gorton, Ian** <i.gorton@northeastern.edu>

To:  Xianhe Zhang

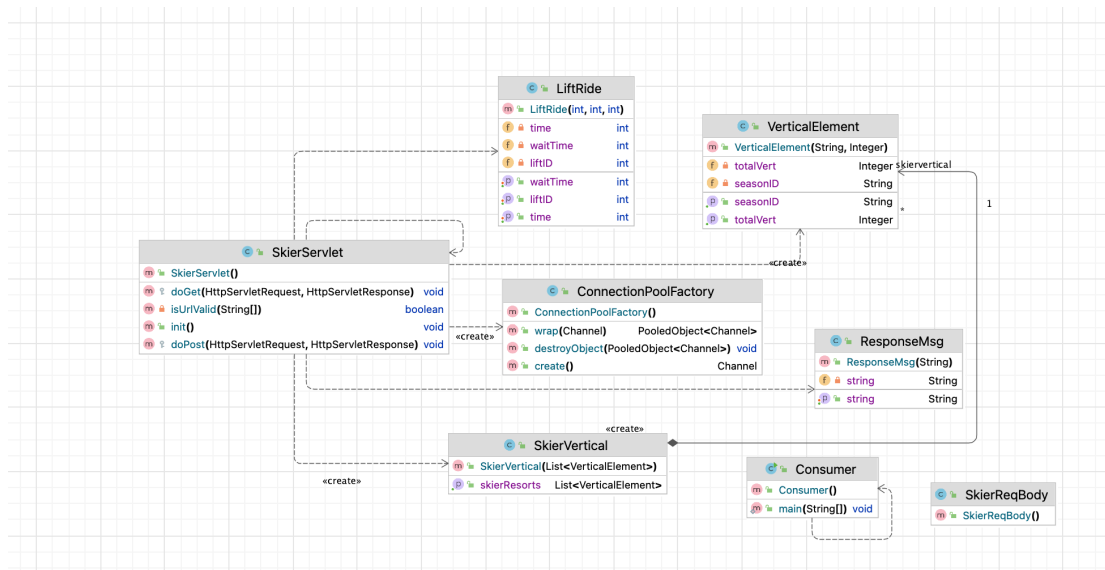
All good. Asap please

Sent via the Samsung Galaxy S21 FE 5G, an AT&T 5G smartphone  
Get [Outlook for Android](#)

## Repo URL

<https://github.com/xianhe-zhang/CS6650-22Fall>

## Design



For this assignment, the server design is pretty much like the assignment 1. But I add connection pool to store the connection threads, this will save resources.

For each thread, we need to take one cached thread from our pool and try to connect to our RabbitMQ (the configuration is below). And our consumer is monitoring specific queues.

```

factory.setHost("35.90.118.182");
factory.setPort(5672);
factory.setUsername("maXXX");
factory.setPassword("maXXXXXXX");
  
```

## Load Balancer

For this assignment, I tried to set up single-instances, double-instances, four-instances.

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
<input type="checkbox"/>	6650-backup	i-0a3a70212625ed5a6	Running	t2.micro	2/2 checks passed	No alarms	us-west-2
<input type="checkbox"/>	6650-consumer	i-00f62d2741abba729	Running	t2.micro	2/2 checks passed	No alarms	us-west-2
<input type="checkbox"/>	6650-server-b2	i-0b267e0874e2af721	Running	t2.micro	2/2 checks passed	No alarms	us-west-2
<input type="checkbox"/>	6650-server-b1	i-0eef3d5cbc22aa313	Running	t2.micro	2/2 checks passed	No alarms	us-west-2
<input type="checkbox"/>	6650-server-b3	i-0b238ddede69adf9c	Running	t2.micro	2/2 checks passed	No alarms	us-west-2

# 6650-target-group

Details

arn:aws:elasticloadbalancing:us-west-2:065157866614:targetgroup/6650-target-group/5ce6c0372911fd81

Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 8080	HTTP1	<a href="#">vpc-0caaf4ad5e</a>
IP address type	Load balancer		
IPv4	<a href="#">6650ServerLB1</a>		

Total targets

Healthy

Unhealthy

Unused

Initial

1

1

0

0

0

Targets

Monitoring

Health checks

Attributes

Tags

Registered targets (1)

Deregister

Filter resources by property or value

<input type="checkbox"/>	Instance ID	Name	Port	Zone	Health status	Health status details
<input type="checkbox"/>	<a href="#">i-0a3a70212625ed5a6</a>	6650-backup	8080	us-west-2c	healthy	

Filter by tags and attributes or search by keyword

1 to 1 of

<input checked="" type="checkbox"/>	Name	DNS name	State	VPC ID	Availability Zones	Type
<input checked="" type="checkbox"/>	6650ServerLB1	6650ServerLB1-917956156....	Active	vpc-0caaf4ad5ecd68eca	us-west-2d, us-west-2c	application

Basic Configuration

Name

6650ServerLB1

ARN

arn:aws:elasticloadbalancing:us-west-2:065157866614:loadbalancer/app/6650ServerLB1/457d5d072c3788a8

DNS name

6650ServerLB1-917956156.us-west-2.elb.amazonaws.com (A Record)

State

Active

Type

application

Scheme

internet-facing

IP address type

ipv4

Edit IP address type

VPC

[vpc-0caaf4ad5ecd68eca](#)

Availability Zones

[subnet-037a4520008ad1b2e - us-west-2d](#)  
IPv4 address: Assigned by AWS

[subnet-0fa8887072de38ff8 - us-west-2c](#)  
IPv4 address: Assigned by AWS

Edit subnets

Created time

2024-08-08 14:00:00 UTC


# Result Analysis

## SINGLE-INSTANCE

Queued messages last ten minutes ?



Message rates last ten minutes ?



---

phase1 has already completed 32000 requests  
phase2 has already completed 168000 requests


Client Part 1 Result:

-----  
Number of successful requests sent: 200000  
Number of unsuccessful requests: 0  
The total run time(wall time): 138105 milliseconds  
The total throughput per Sec: 1448

Client Part 2 Result:

-----  
200000  
7015849.0  
Mean response time: 35.079245  
Median response time: 31.0  
Throughput: 28.506884911576634  
99th response time: 129.0  
min and max response time: min: 12.0 , max: 727.0  
The total throughput per Sec: 1448  
phase duration: 138105

## TWO-INSTANCES



phase1 has already completed 32000 requests  
phase2 has already completed 168000 requests

### Client Part 1 Result:


-----  
Number of successful requests sent: 200000  
Number of unsuccessful requests: 0  
The total run time(wall time): 113993 milliseconds  
The total throughput per Sec: 1754

### Client Part 2 Result:


-----  
200000  
5726710.0  
Mean response time: 28.63355  
Median response time: 23.0  
Throughput: 34.92406634874125  
99th response time: 158.0  
min and max response time: min: 12.0 , max: 986.0  
The total throughput per Sec: 1754  
phase duration: 113993

## FOUR-INSTANCES

Queued messages last ten minutes ?



Message rates last ten minutes ?



phase1 has already completed 32000 requests  
phase2 has already completed 168000 requests

Client Part 1 Result:

```
-----  
Number of successful requests sent: 200000  
Number of unsuccessful requests: 0  
The total run time(wall time): 102000 milliseconds  
The total throughput per Sec: 1960
```

Client Part 2 Result:

```
-----  
200000  
4999586.0  
Mean response time: 24.99793  
Median response time: 20.0  
Throughput: 40.00331227425631  
99th response time: 141.0  
min and max response time: min: 11.0 , max: 931.0  
The total throughput per Sec: 1960  
phase duration: 102000
```

For SERVER, I still have 2 phases like Assignment 1. For phase 2, I use **54** threads to send the rest requests.

For CONSUMER, **basicQos == 30**, because publisher rate is way higher than consumer rate

From the performance screenshot above, we can tell load balancer is good for our throughput, given more instances are backing up, throughput improves accordingly.